

CLIENT BRIEFING

The operations assistant that's **allowed to act**.

A secure, auditable AI assistant that moves Acme's sales, support and operations teams from fragmented manual investigation to governed, AI-assisted action — without ever handing the keys to the model.

—— The LLM **advises and explains**. Deterministic policy **executes**.

Customer context is scattered — and the AI that touches it can't be trusted to act.

01 • TODAY

Fragmented manual investigation

Staff stitch together the customer profile, open issues, history and the right next step across disconnected systems — slow, inconsistent, and easy to get wrong under SLA pressure.

02 • THE RISK

Generic AI can't be trusted in operations

Off-the-shelf assistants hallucinate, can't prove their claims, and — given tools — may act without authority, inventing actions or writing to systems they shouldn't touch.

03 • THE GAP

No audit trail when AI decides

When a model influences an operational decision, there's no record of who acted, on what evidence, under what permission, or with what outcome. That's a non-starter for the enterprise.

ONE PRINCIPLE, ENFORCED EVERYWHERE

The LLM **advises**. Deterministic policy **executes**.

01 • EVIDENCE

Evidence-first answers

Every claim links to the exact records and tools that support it. Unsupported answers are flagged "Insufficient Evidence" — never asserted.

02 • CONTROL

Propose → confirm → execute

No side-effecting write happens without an explicit human confirmation. The model can recommend, but a person and a policy gate pull the trigger.

03 • GUARDRAILS

Closed catalogue + server-side RBAC


The model can't invent action types or grant itself rights. Permissions come from the verified identity — never from the model's plan.

04 • AUDIT

An append-only decision ledger

Every action records who acted, why, on what evidence, under what permission, with what outcome — a record that can never be rewritten.

Every request flows through governance gates — and the LLM only ever sits in the middle.



Governed, observable and provable — **by design.**

✓ Grounded ▶ Proposed × Denied

A verification badge on every answer

Grounded, Insufficient Evidence, Permission Denied, Action Proposed — at-a-glance status, never a silent guess.

server-side

RBAC enforced from the token

Authorization is taken from the verified identity, never the model's plan. A plan that says "admin" grants nothing.

append-only

An audit record that can't be rewritten

The Decision Ledger only ever appends. The history of who did what, and why, is permanent and tamper-evident.

\$ / turn

Cost, tokens & latency on every turn

"What does this cost per query?" has a numeric answer at all times — surfaced in the trace, persisted in the ledger.

18 × 3

An evaluation suite with variance

18 cases run three times across tool selection, grounding, RBAC, action reasonableness and adversarial defence.

swappable

Provider-agnostic, ready to integrate

Anthropic, OpenAI, Google or local models behind one interface; the same tool contracts can wrap Salesforce, ServiceNow or SAP.